University of Minho
Department of Informatics

LABORATÓRIO DE INSTRUMENTAÇÃO E FÍSICA EXPERIMENTAL DE PARTICULAS
LIP

EFFICIENT PROCESSING OF ATLAS EVENTS ANALYSIS IN HOMOGENEOUS AND HETEROGENEOUS PLATFORMS WITH ACCELERATOR DEVICES

André Pereira
Prof. Alberto Proença (Advisor)
Prof. António Onofre (Co-Advisor)

Overview

- Motivation
- Structure of `ttH_dilep`
  - Analysis of critical regions
- Improving efficiency through parallelism
  - On shared memory homogeneous systems
  - On heterogeneous systems with a GPU accelerator
  - On heterogeneous systems with a MIC accelerator
  - Efficient data & workload scheduling

Motivation

- The target model for an event

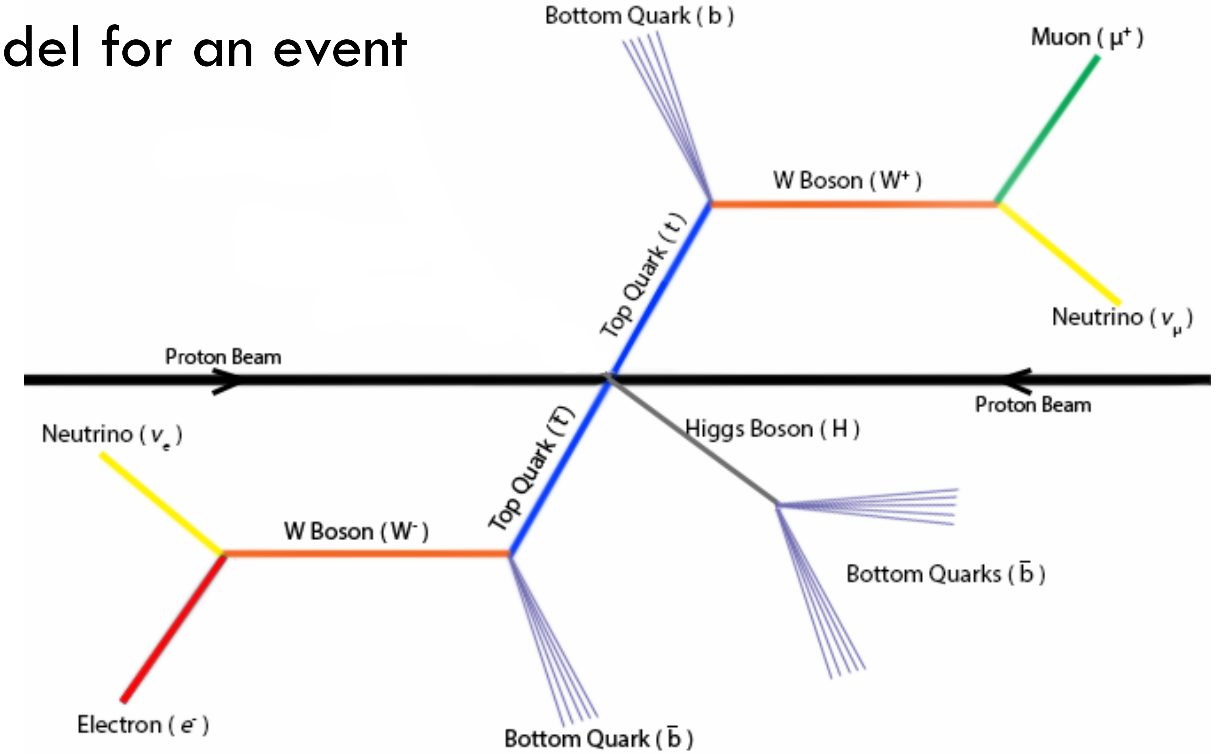

- Reconstruction of the Top Quarks (t&t̄) system with Higgs boson
  - the analysis and reconstruction code => ttH_dilep
- Goal: to improve the efficiency of ttH_dilep

Structure of ttH_dilep (3)

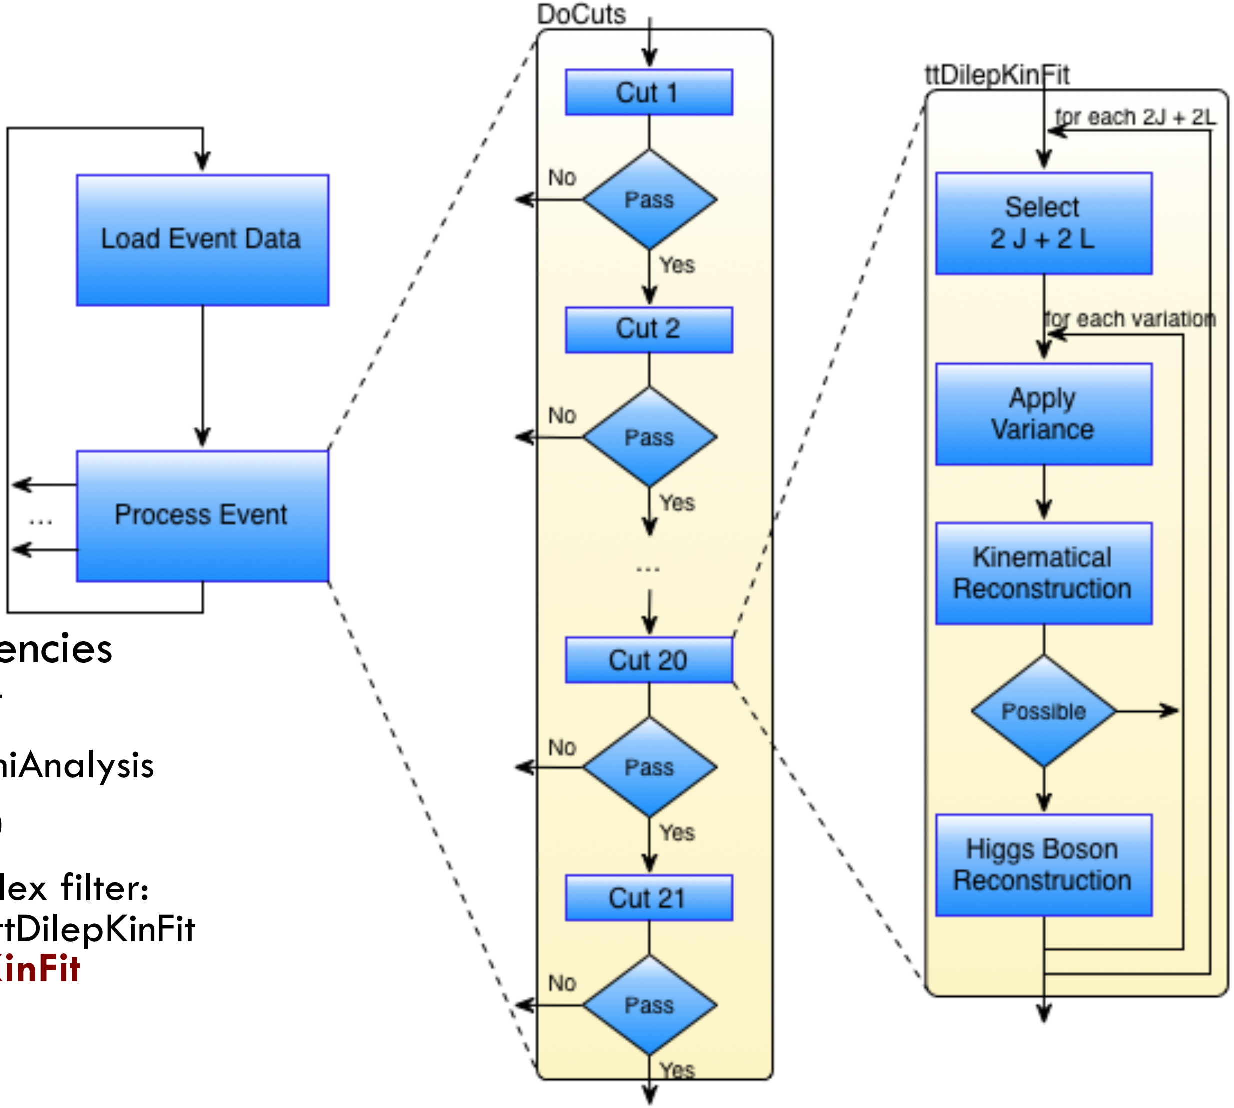

- Dependencies
  - ROOT
  - LipMiniAnalysis
- Cut #20
  - Complex filter:
    ttDilepKinFit
    aka **KinFit**

Critical regions

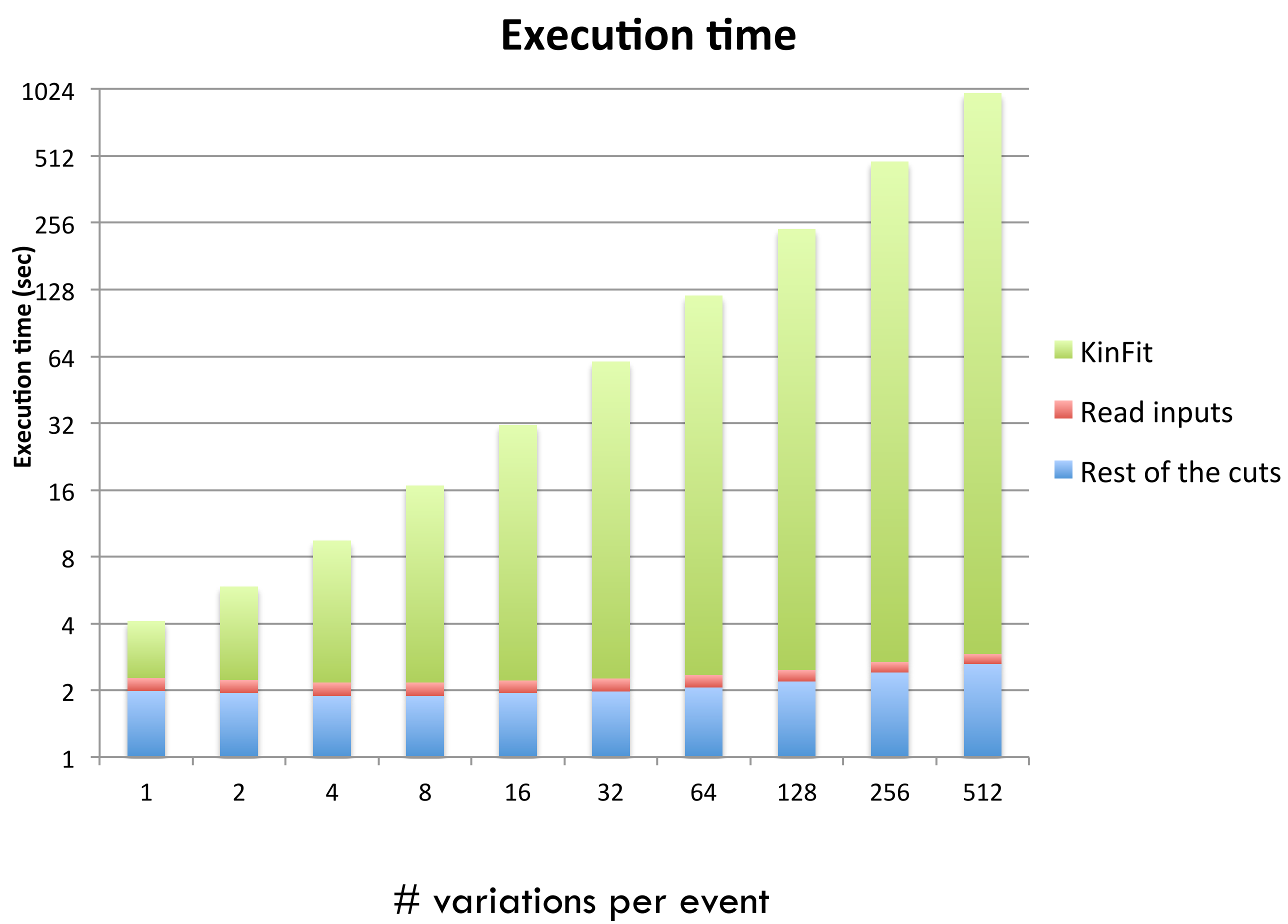

Improving efficiency with parallelism…

Sequential

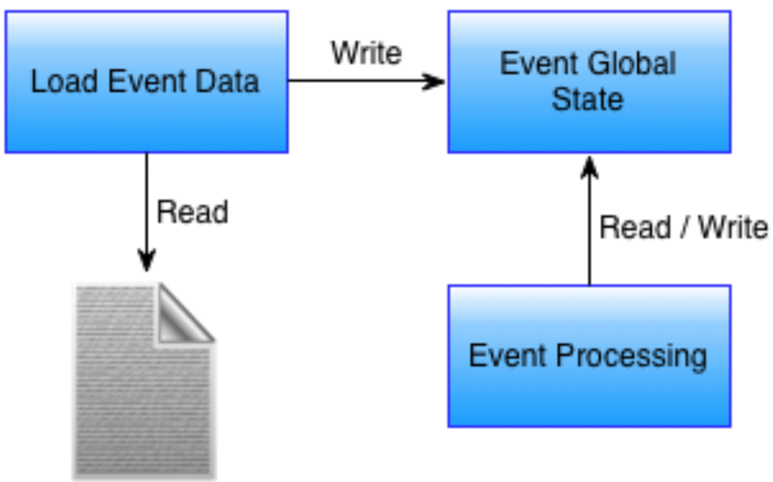

Parallel

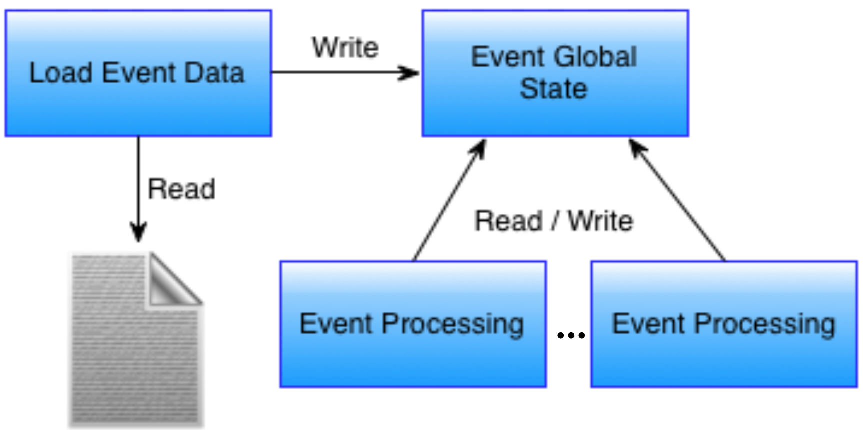

… no way with single global state!

Sequential

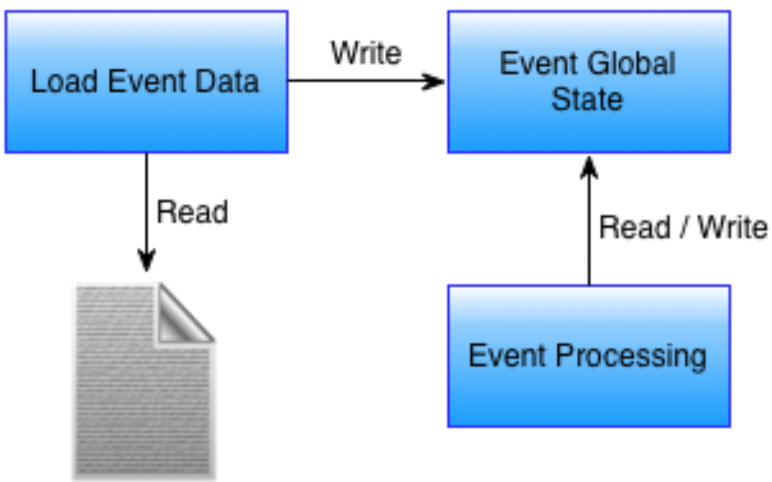

Parallel

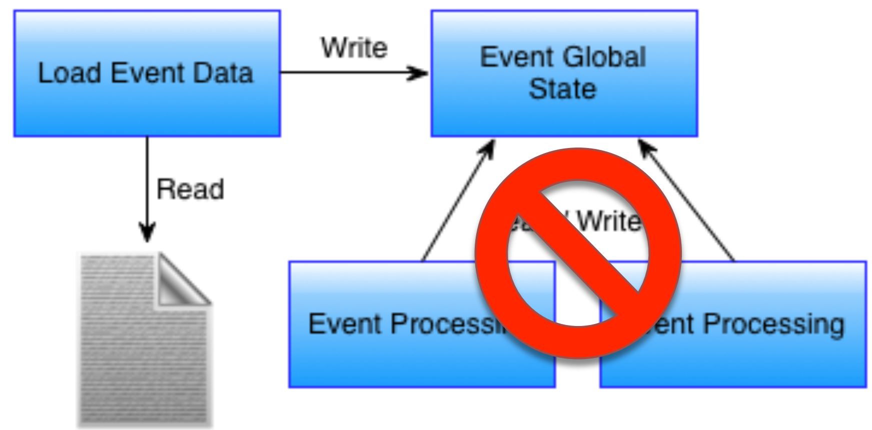

Approach 1: parallelize KinFit, shared memory, no h/w accelerators

Sequential

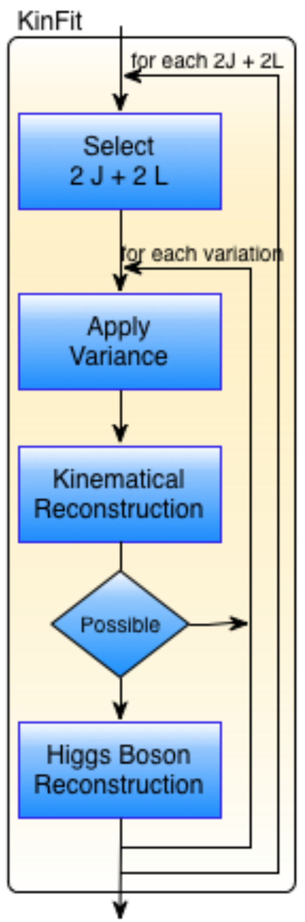

Parallel

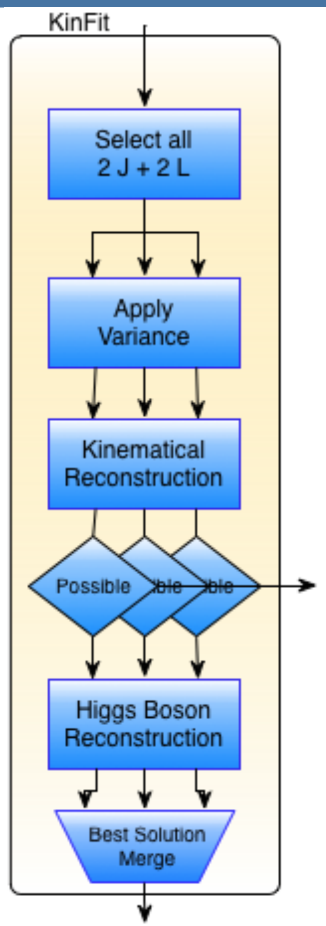

- Selects all sets of 2 jets & 2 leptons, and builds a new data structure for each set
- In parallel:
  - applies a tolerance (variation) to each measure
  - performs a kinematical and Higgs boson reconstructions
  - selects the best solution (reduction)
- The new data structure:
  - all data are replicated; or
  - common data share a pointer

Performance analysis (1)

System: 2 x Intel E5-2670 (total 16 cores, 32 threads)

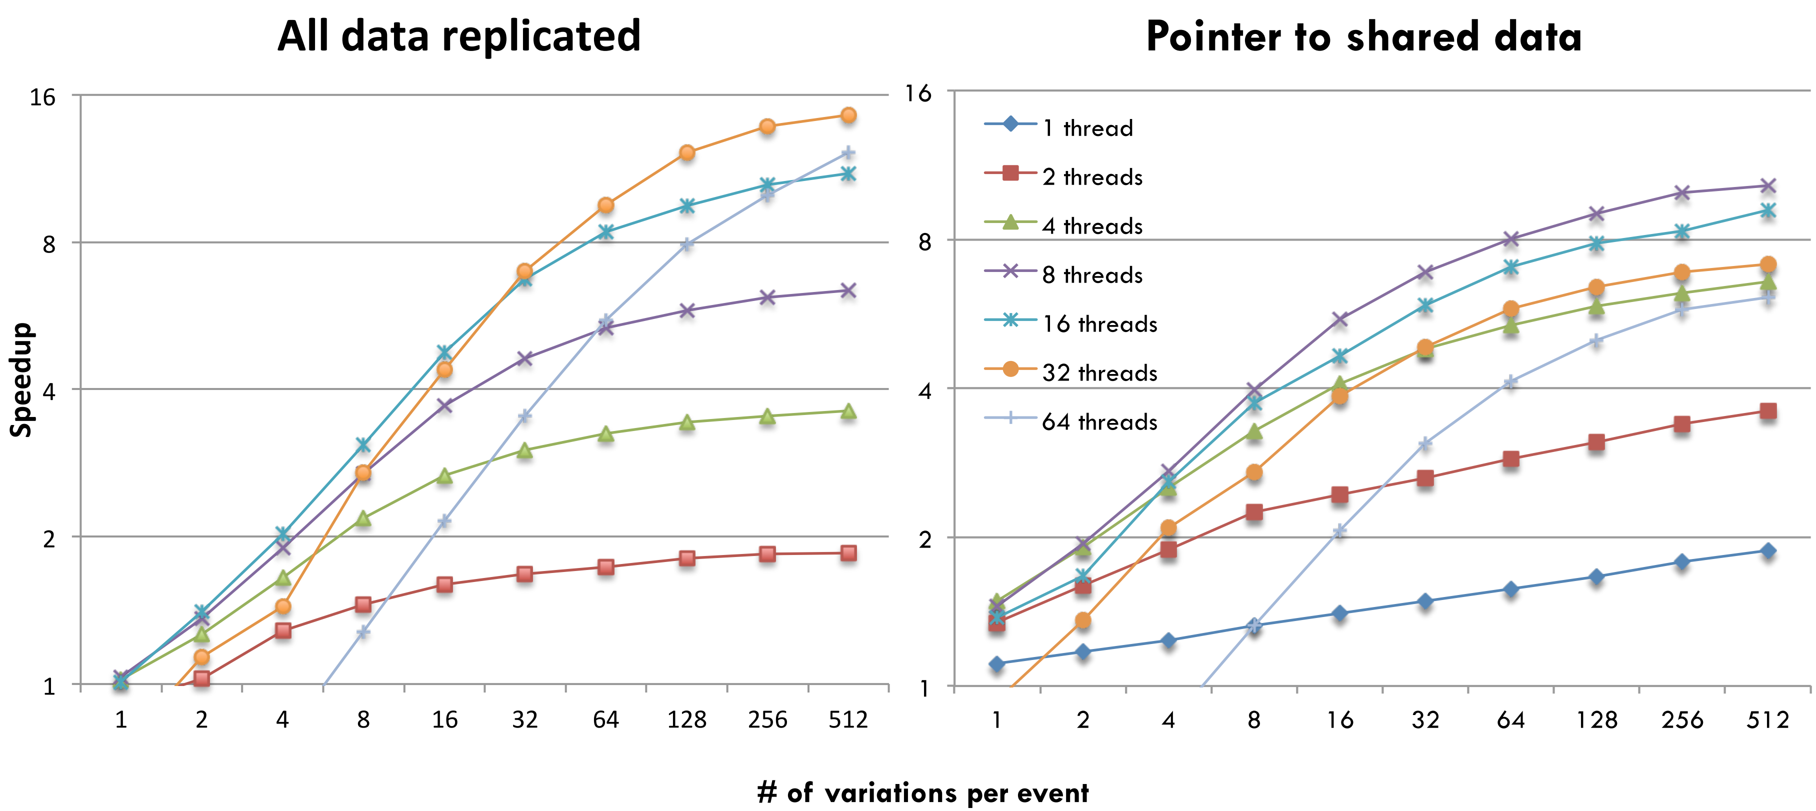

Plots: speedup *versus* the original sequential version

Approach 2: parallelize KinFit, with GPU accelerator (distributed memory)

Parallel, no accelerator

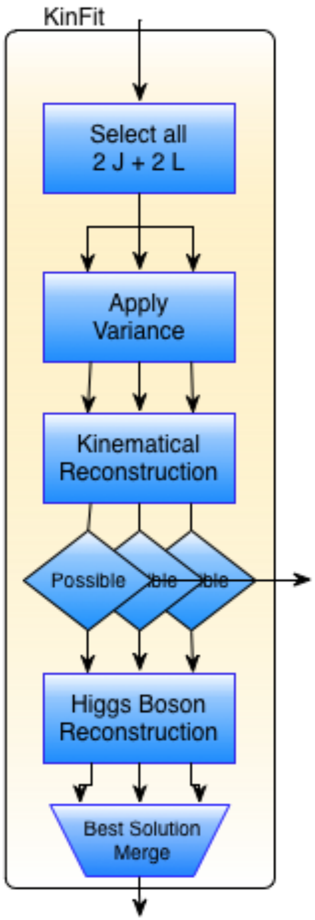

Parallel: multicore + GPU

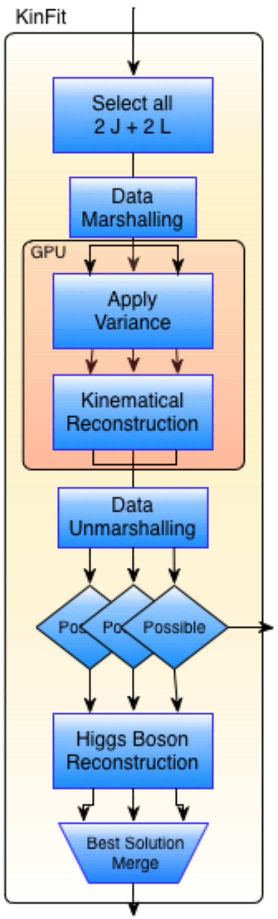

- Data (un)marshalling:
  - Transform ROOT (and application classes) in arrays to transfer to and from the GPU
- In the GPU:
  - Apply the variance on the inputs
  - Perform the kinematical reconstruction
- Parallel reconstruction of the Higgs boson on CPU

Performance analysis (1)

System: 2 x AMD Opteron 6174 (total 24 cores) and NVidia Tesla Fermi C2050 GPU

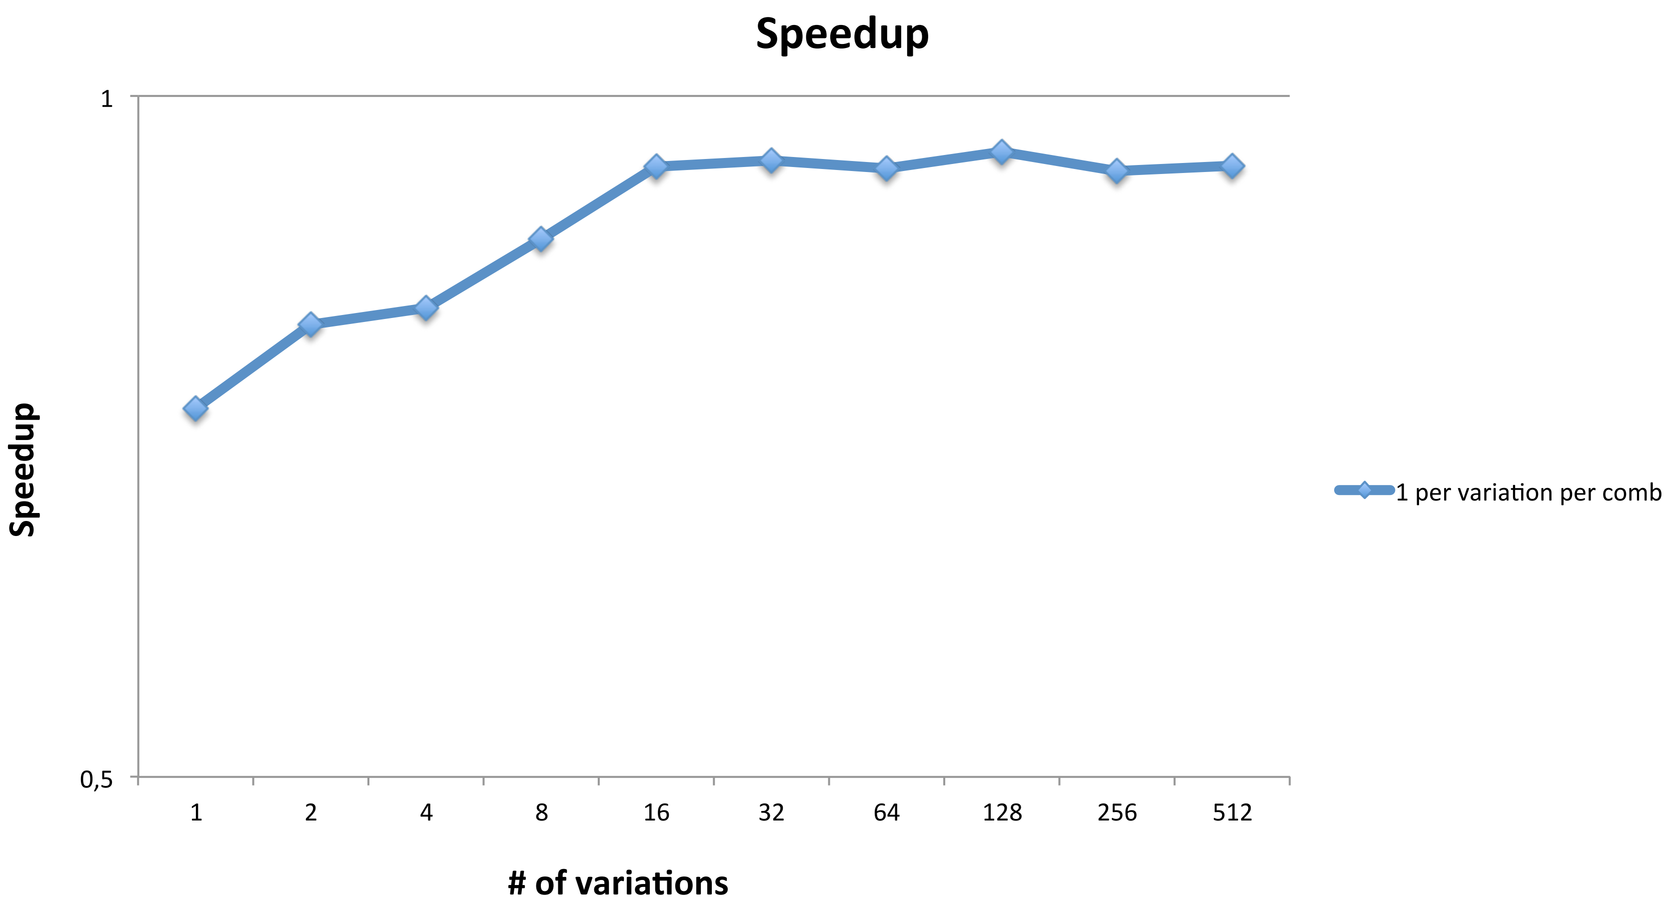

Performance analysis (2)

Ideal execution

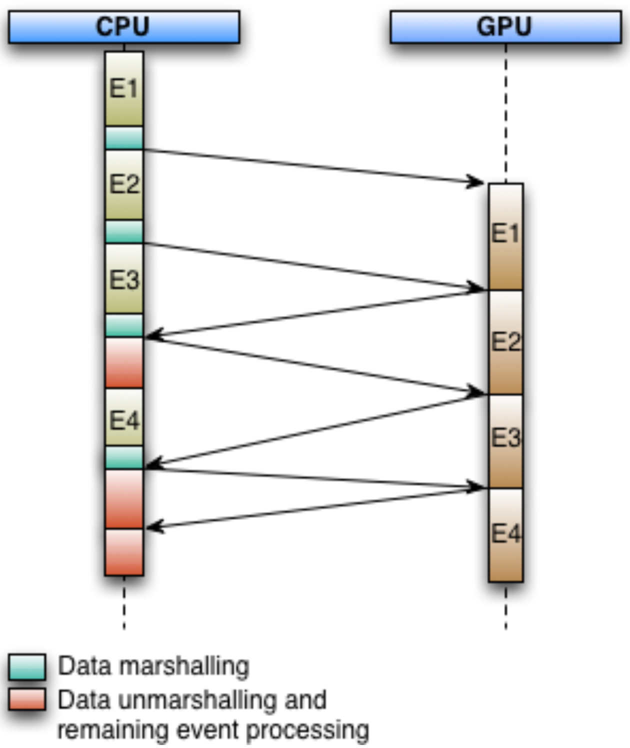

Current execution

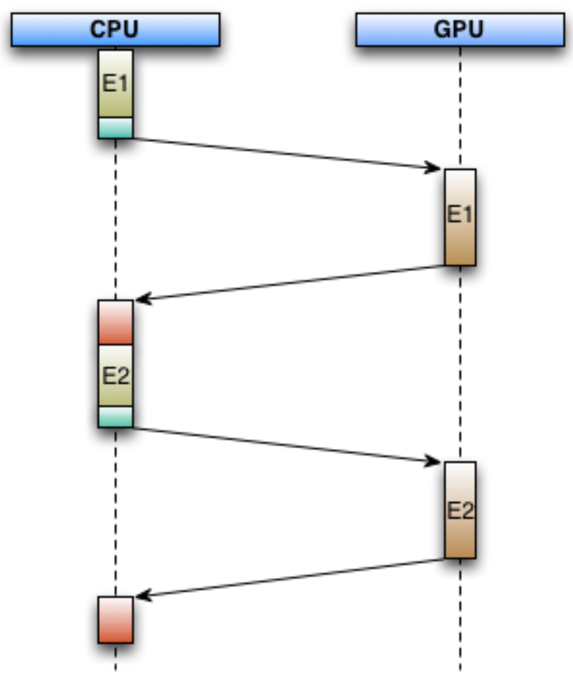

For 256 variations the CPU is idle 31% of the time…

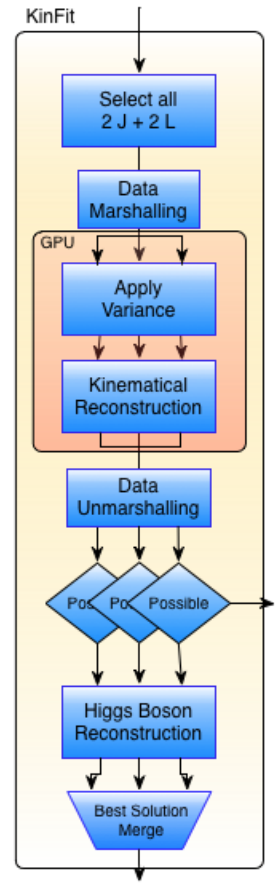

A single global state **prevents** simultaneous CPU & GPU processing

Approach 3: events from different files

Sequential

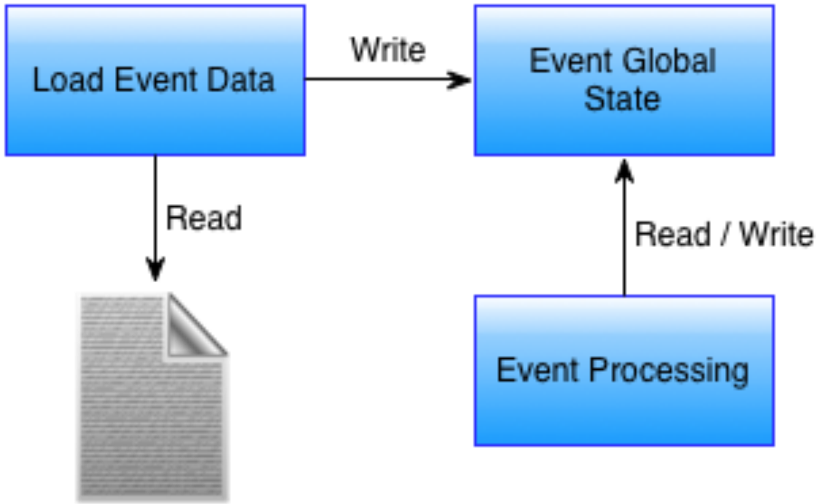

Parallel

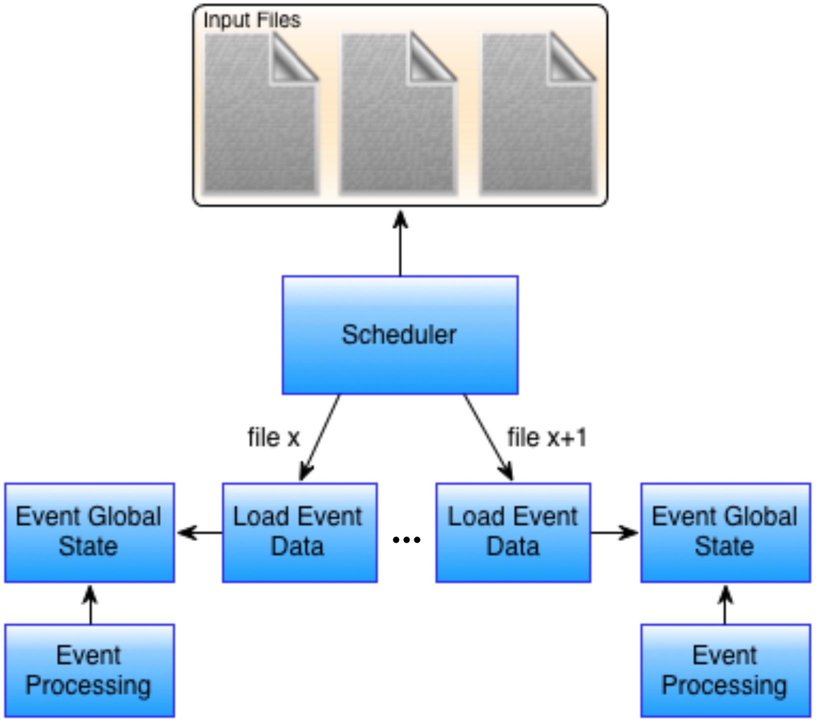

Performance analysis (1)

System: 2x Intel E5-2670 (total 16 cores, 32 threads)

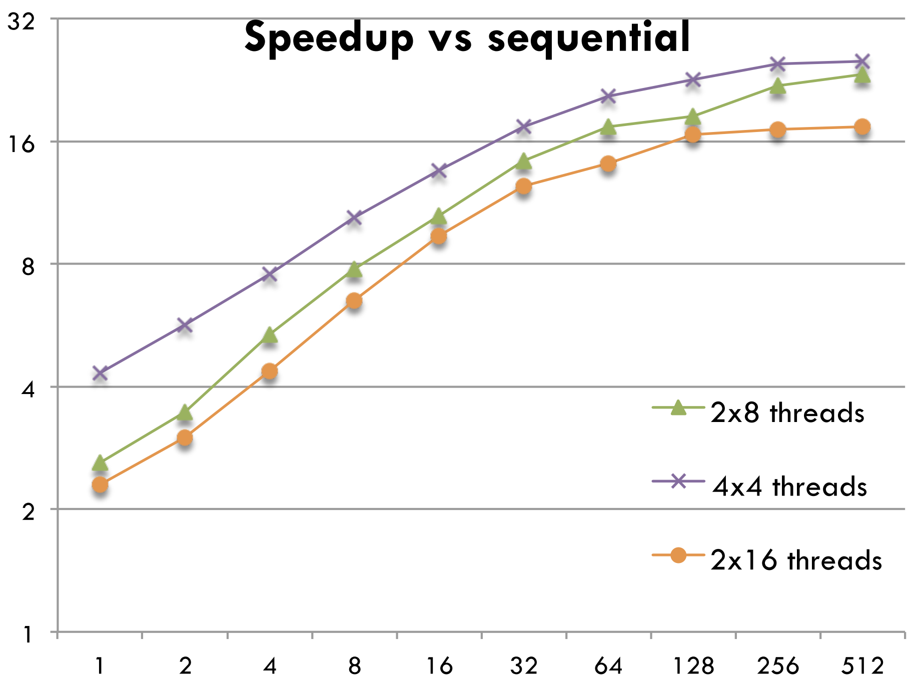

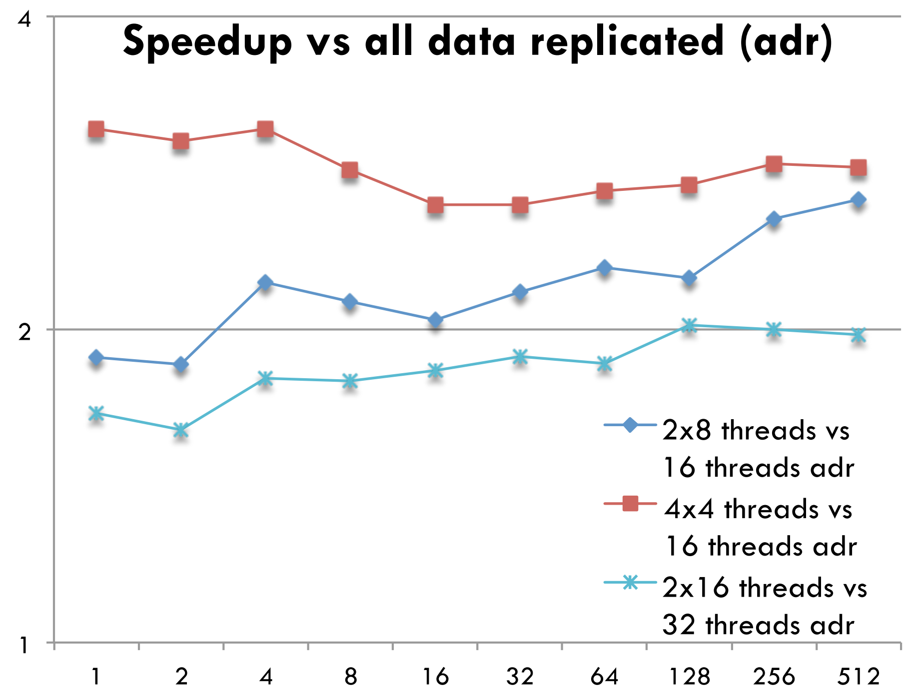

# of variations

Conclusions

- Homogeneous systems
  - All data replicated provides the best performance (multiple CPUs)
  - Pointer to shared data provides the best efficiency (single CPU)
- Heterogeneous systems
  - Unefficient GPU usage due to the lack of a global data structure
  - Xeon Phi preliminary implementation limited by the drivers
- Scheduler provides the best efficiency for multi-CPU systems
  - Allied to a parallelization of LipMiniAnalysis a tool can be developed that automatically extracts parallelism

University of Minho
Department of Informatics

LABORATÓRIO DE INSTRUMENTAÇÃO E
FÍSICA EXPERIMENTAL DE PARTICULAS

EFFICIENT PROCESSING OF ATLAS EVENTS ANALYSIS IN HOMOGENEOUS AND HETEROGENEOUS PLATFORMS WITH ACCELERATOR DEVICES

André Pereira
Prof. Alberto Proença (Advisor)
Prof. António Onofre (Co-Advisor)